ModelDoctor

Diagnosing Silent Failures in Machine Learning Systems

Presented by Team ASTROID
Team Member 1 – Prabhas N
Team Member 2 – Poornima Bhat

Problem Statement

In traditional IT, a bug causes a crash. In Machine Learning, failures are silent. The code runs, the GPU burns power, and training completes—but the model learns the wrong thing.

The Silent Killers: Issues like Data Leakage, Gradient Instability, and Overfitting are invisible to standard compilers.

The Industry Crisis: Developers waste 40% of their productivity on manual "guesswork" debugging, and companies lose thousands of dollars in Compute Hemorrhage (wasted GPU hours) on models that were doomed from the start.

The Diagnostic Gap: We have "Unit Testing" for software, but we have no "MRI" for ML logic.

The Silent Killer: Undetected ML Failures

Machine learning models often fail in a way that is insidious and difficult to detect: silently. Unlike traditional software that crashes, an ML model can appear to function perfectly on the surface, only to betray expectations in the real world.

Code Runs Successfully

All lines of code execute without exceptions.

Training Completes

The model finishes its training epochs, reporting no errors.

Accuracy Looks Good

Initial performance metrics seem promising, sometimes even excellent.

Despite these apparent successes, the model's real-world performance falls short. Critically, there's no clear warning or immediate explanation for this discrepancy.

Introducing ModelDoctor: Your ML Diagnostic Expert

ModelDoctor is a revolutionary diagnostic system designed to proactively identify and explain silent failures within your machine learning pipelines. We treat your ML pipeline with the same methodical approach a doctor would use for a patient, focusing on deep root-cause analysis.

Treats ML like a Patient

Systematic examination of your pipeline's health.

Root-Cause Diagnosis

Pinpoints the fundamental issues, not just symptoms.

Clear Explanations

Articulates what is wrong, why it's happening, and its implications.

Silent ML failures range from subtle to severe impact.

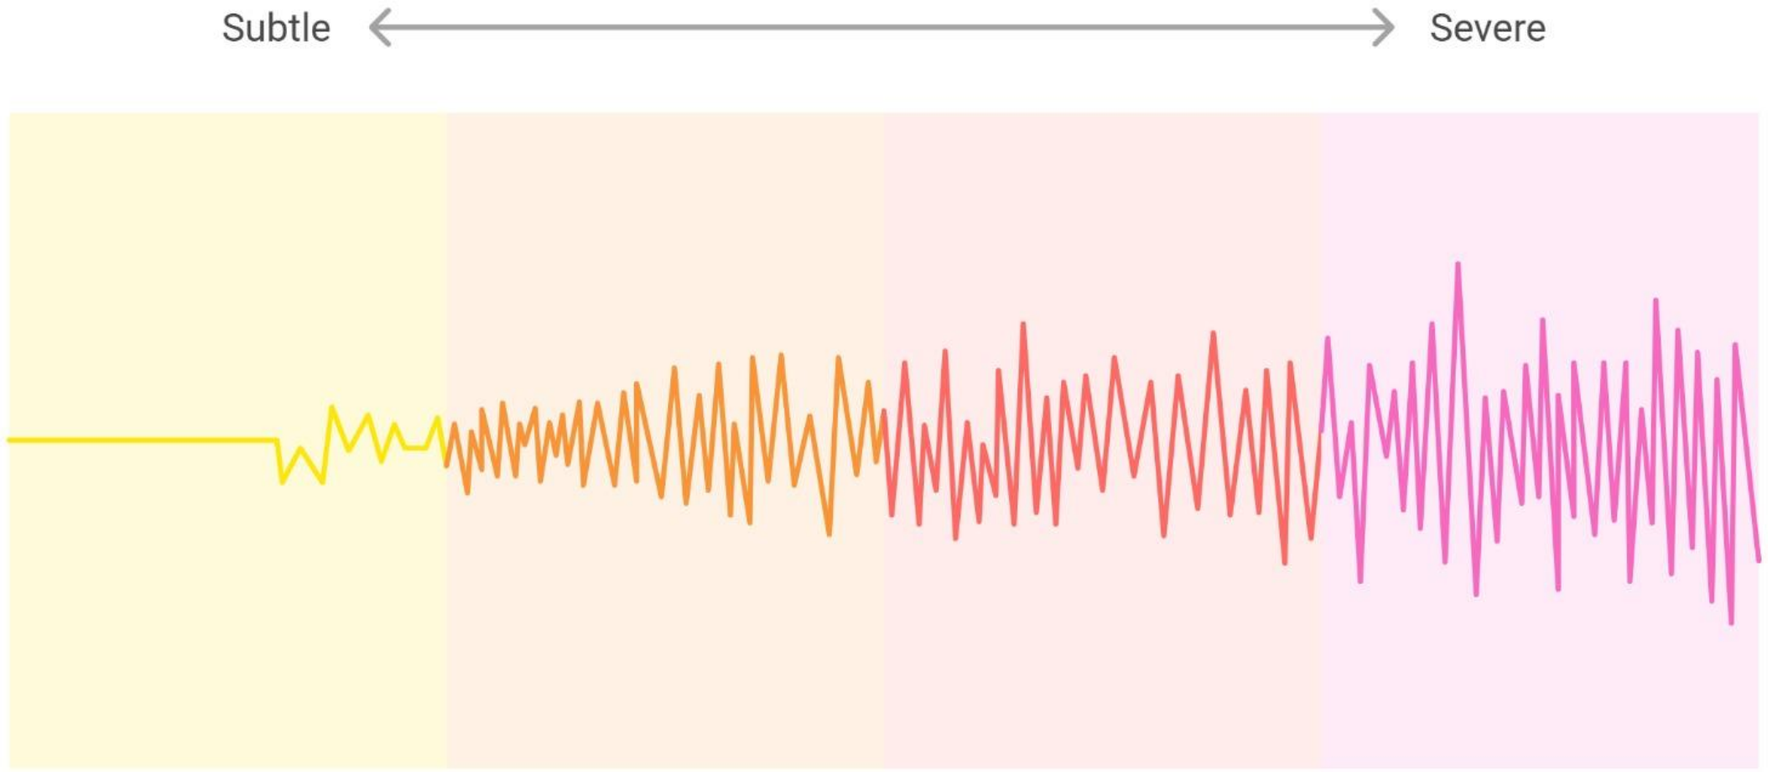

Data Leakage

Inadvertently influences training

Overfitting

Performs well on training, poorly on unseen

Incorrect Split

Biased partitioning leads to unreliable models

Feature Misuse

Uses unavailable features during inference

ModelDoctor Diagnostic Workflow

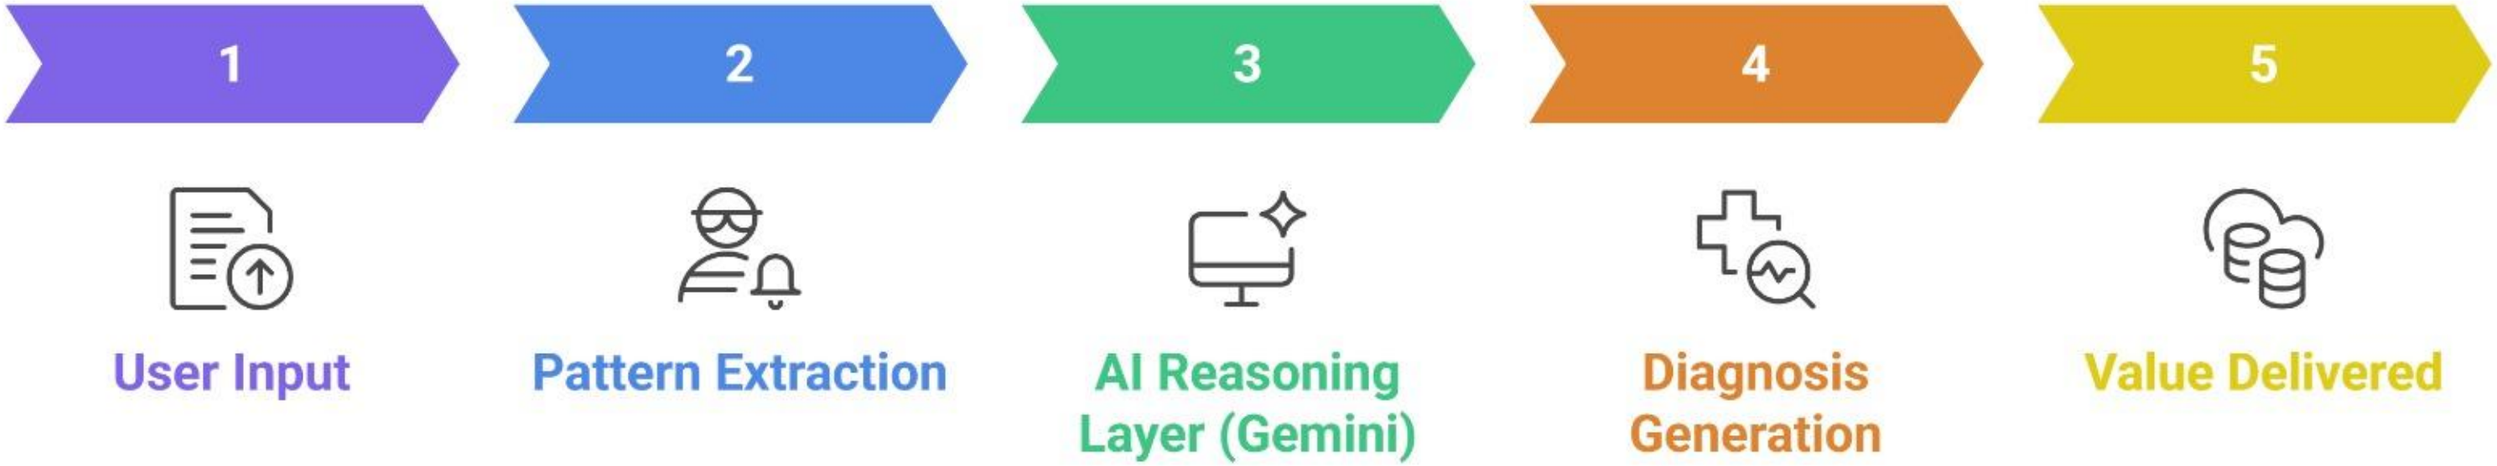

Leveraging Gemini API for Advanced Reasoning

The Gemini API is at the heart of ModelDoctor's diagnostic capabilities, enabling sophisticated analysis and human-like reasoning to uncover complex ML issues.

Contextual Understanding

Reads vast amounts of code and log context, understanding relationships across diverse data types.

Pattern Recognition

Connects intricate patterns between inputs that indicate hidden problems.

Reasoning About Failures

Reasons about the underlying causes of failures, providing deep insights rather than superficial observations.

Actionable Explanations

Generates clear, human-readable explanations and actionable fixes for detected issues.

ML Development Lifecycle Improvement

Inefficient Development

Slow, costly, unreliable ML models

Early Problem Detection

Identify issues before they escalate

Data Cleaning & Bias Prevention

Ensure clean, unbiased data for training

Model Learning Prevention

Prevent incorrect model learning

Efficient & Trustworthy ML

Reliable, robust, and cost-effective models

AI-Powered Model Diagnosis Pipeline

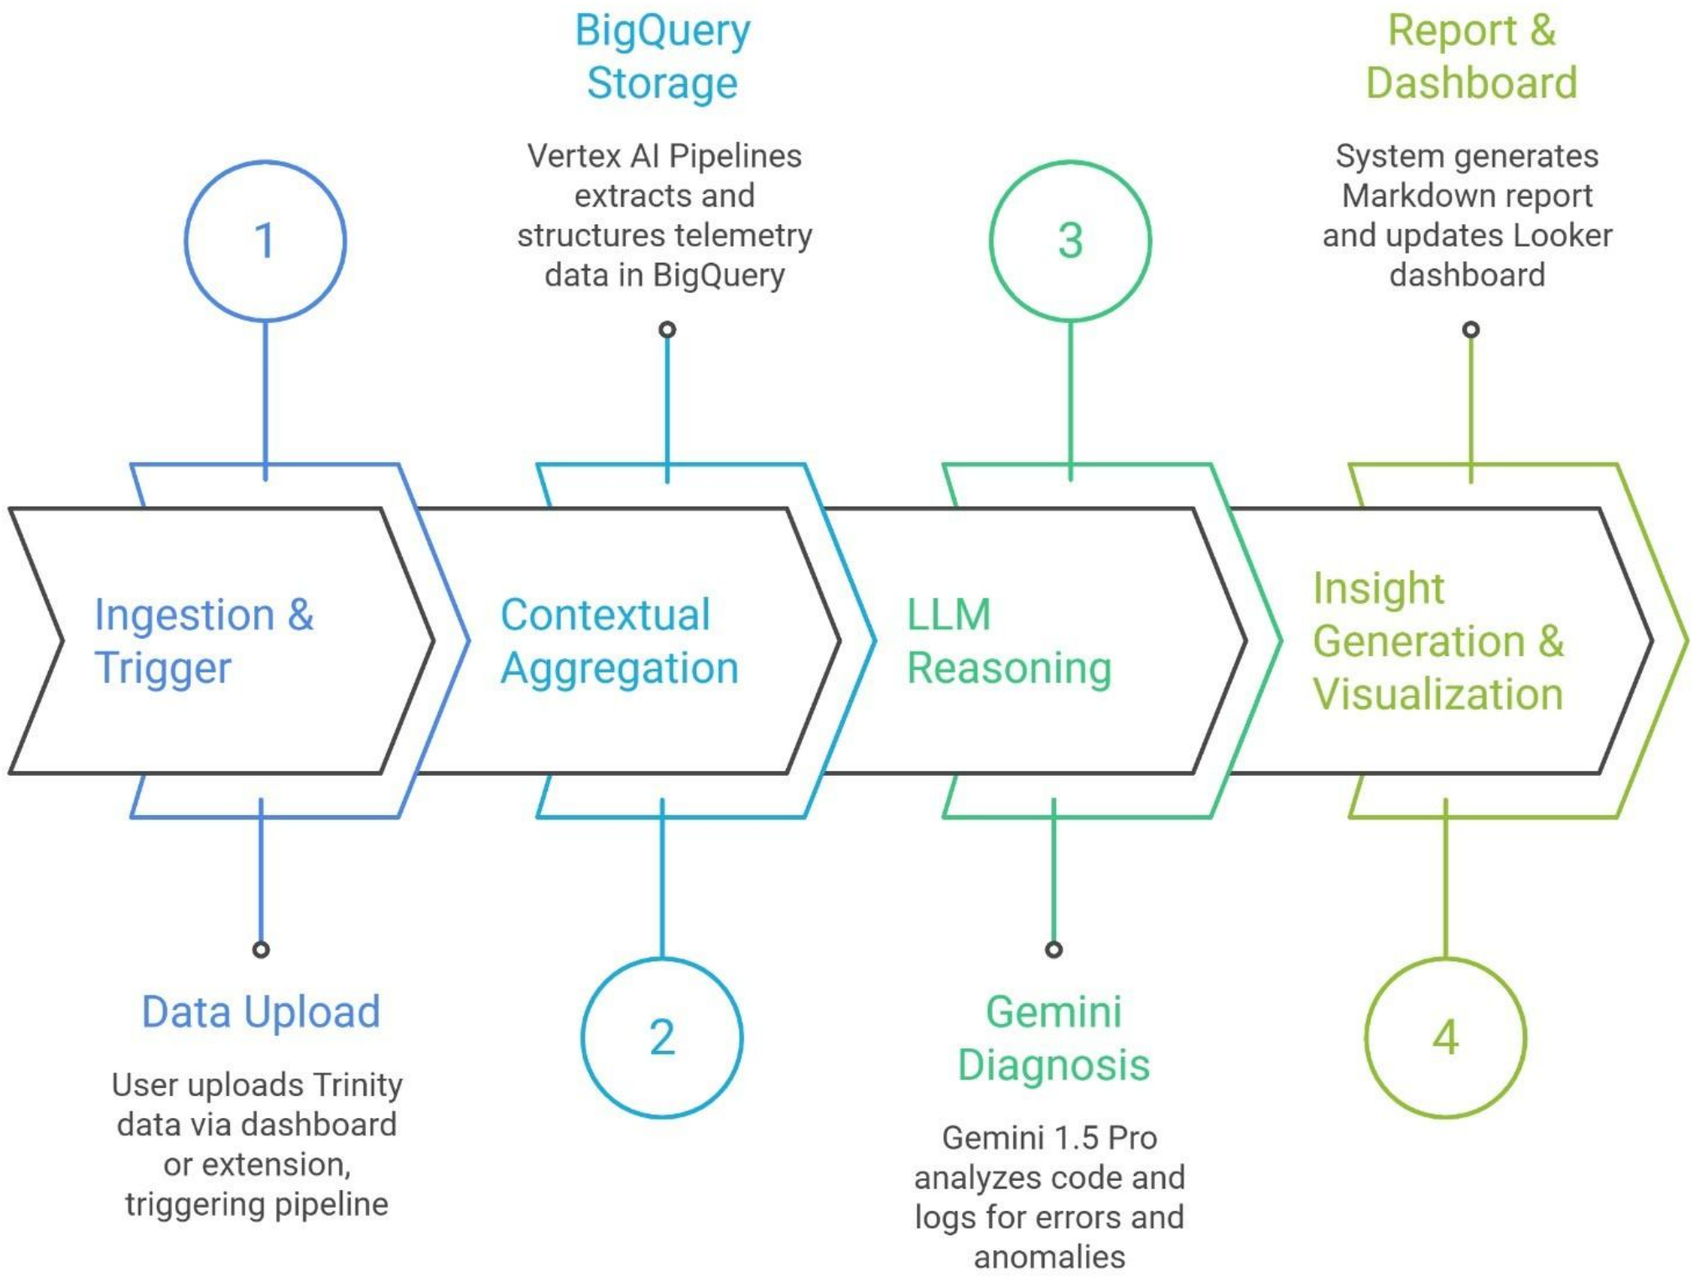

Comprehensive Output: Your Diagnosis Report

The ModelDoctor diagnosis report is designed to be clear, concise, and most importantly, actionable. Each report provides all the necessary information to understand and rectify detected issues efficiently.

Detected Issue

A precise description of the identified silent failure.

Severity Level

Categorisation of the issue's impact (e.g., Critical, High, Medium, Low).

Explanation of Cause

Detailed breakdown of why the issue occurs and its potential consequences.

Suggested Fix

Recommendations on how to resolve the identified problem.

Location for Application

Guidance on where in the code or pipeline the fix should be applied.

Conclusion: Trustworthy ML from the Start

ModelDoctor addresses a critical gap in machine learning development by proactively detecting silent failures, fostering a new era of trust and reliability in AI systems.

Detects Silent ML Failures Early

By integrating and analysing all relevant signals, ModelDoctor catches hidden issues before they escalate, saving time and resources.

Encourages Reasoning-Driven Debugging

Provides clear explanations and actionable fixes, moving away from guesswork to intelligent problem-solving.

With ModelDoctor, you can build, test, and deploy machine learning models with confidence, knowing that potential silent failures have been identified and addressed, ensuring reliable AI before it reaches production.